

1 Question: How would you find the boundary of the diffusion of each dye? (Note: If you have  
multiple answers in mind, break them apart and explain each one separately.) Explain each  
solution/algorithm in detail.

2

3 Answers: The following steps are used to find diffusion boundaries of each dyes:

4 1. Extract color channels corresponding to each dyes (Red: channel 0, Green:  
channel 1).

5 2. Convert single-channel images to binary images using thresholding (Red: 0.06,  
Green: 0.06)

6 3. Connect neighboring pixels by filling small holes less than a certain pixels  
(Red: 64 pixels, Green: 128 pixels).

7 4. Remove noises by eliminating connected components that are less than a certain  
threshold (Red: 20000 pixels, Green: 64 pixels).

8 5. Soften edges by dilating the image with a certain matrix (Red: 3x3, Green: 4x4).

9 6. Find edges using Roberts Cross filter.

```
1 from skimage.io import imread
2 import matplotlib.pyplot as plt
3 from skimage.filters import roberts
4 from skimage.morphology import (
5     dilation,
6     square,
7     remove_small_objects,
8     remove_small_holes,
9 )
10 import numpy as np
11 from pathlib import Path
12
13 ROOT = Path(__file__).parent
14
15 def find_red_diffusion_boundary(img: np.ndarray) -> np.ndarray:
16     """Find the red diffusion boundary.
17
18     Args:
19         img (np.ndarray): images.
20
21     Returns:
22         np.ndarray: binary images where 1s indicate the boundary.
23     """
24     # Extract the red channel from the image and normalize it.
25     img = img[:, :, 0] / 255.0
26
27     # Convert the image to binary image with thresholding.
28     img = img > 0.06
29
30     # Connected neighboring pixels by filling small holes that are less than 64 pixels.
31     img = remove_small_holes(img, 64)
32
33     # Reduce noises by eliminating connected components that are less than 20,000 pixels.
34     img = remove_small_objects(img, 20000)
35
36     # Soften edges by dilating the image with 3x3 matrix
37     img = dilation(img, square(3))
38
39     # Fill in the remaining holes
40     img = remove_small_holes(img, 128)
41
42     # Use roberts filter to find edges
43     img = roberts(img)
44     img = img != 0
45
46     return img
47
48
49 def find_green_diffusion_boundary(img: np.ndarray) -> np.ndarray:
50     """Find the green diffusion boundary.
51
52     Args:
53         img (np.ndarray): images.
54
55     Returns:
56         np.ndarray: binary images where 1s indicate the boundary.
57     """
```

```
58 # Extract the green channel from the image and normalize it.
59 img = img[:, :, 1] / 255.0
60
61 # Convert the image to binary image with thresholding.
62 img = img > 0.06
63
64 # Connected neighboring pixels by filling small holes that are less than 128 pixels.
65 img = remove_small_holes(img, 128)
66
67 # Reduce noises by eliminating connected components that are less than 64 pixels.
68 img = remove_small_objects(img, 64)
69
70 # Soften edges by dilating the image with 4x4 matrix
71 img = dilation(img, square(4))
72
73 # Fill in the remaining holes
74 img = remove_small_holes(img, 1024)
75
76 # Use roberts filter to find edges
77 img = roberts(img)
78 img = img != 0
79
80 return img
81
82
83 def problem_1(in_img: str = ROOT/"dyes.png", out_img: str = ROOT/"problem_1_result.png") ->
None:
84     """Find diffusion boundaries of the red and the green pixels.
85
86     Args:
87         in_img (str, optional): path to the input image. Defaults to "./dyes.png".
88         out_img (str, optional): path to the output image. Defaults to
89     """
90     # Load the image
91     img = imread(in_img)
92
93     # Find boundaries
94     red_diffusion_boundary = find_red_diffusion_boundary(img)
95     green_diffusion_boundary = find_green_diffusion_boundary(img)
96
97     # Find boundaries' coordinates
98     red_coordinates = np.where(red_diffusion_boundary == 1)
99     green_coordinates = np.where(green_diffusion_boundary == 1)
100
101     # Plot boundaries
102     img[red_coordinates[0], red_coordinates[1]] = [255, 0, 0]
103     img[green_coordinates[0], green_coordinates[1]] = [0, 255, 0]
104
105     # Save the resulting image
106     plt.imsave(out_img, img)
107
108
109 if __name__ == "__main__":
110     problem_1()
```

1 Question: How would you create your model to classify the digits? (Note: If you have multiple answers in mind, break them apart and explain each one separately.) Explain each solution/algorithm in detail.

2  
3 Answers: The following steps are used to create a model that is able to recognize handwritten digits.

- 4 1. Load data and labels from mat files (data: 784x2500, labels: 1x2500).
- 5 2. Reshape data and labels (data:2500x784, labels:2500)
- 6 3. Perform K-Fold Cross Validation to split the dataset into 15 folds.
- 7 4. For each fold,...
- 8 4.1. Split the train dataset into a train dataset and a validation dataset.
- 9 Keep the test dataset the same.
- 10 4.2. Load all three datasets into Dataset classes to convert all values to Tensor.
- 11 4.3. Load train\_dataset into a Dataloader class to perform batch training.
- 12 4.4. Define the model, the loss function, and the optimizer.
- 13 4.5. At each epoch, ...
- 14 4.5.1. Feed data from train\_dataloader into the model to produce predicted labels.
- 15 4.5.2. Use the loss function to calculate a loss between predicted labels and truth labels.
- 16 4.5.3. Calculate gradients of each model's weights with respect to the loss.
- 17 4.5.4. Use the optimizer to update a model's weights.
- 18 4.5.5. Calculate train\_loss and train\_accuracy.
- 19 4.5.6. Validate the model with val\_dataset and calculate val\_loss and val\_accuracy.
- 20 4.5.7. If val\_loss has not decreased in the last 10 epochs, end the training process.
- 21 4.5.8. Else, train the model until the final epoch.
- 22 4.6. Test the model with test\_dataset and calculate test\_loss and test\_accuracy.
- 23 5. Calculate the average loss and the average accuracy for all three datasets across all folds.

24 Design details:

- 25 • Model: Fully-connected Model
- 26 Given: in\_features, num\_classes
- 27 Model(
- 28 Linear(in\_features=in\_features, out\_features=32)
- 29 ReLU()
- 30 Linear(in\_features=32, out\_features=32)
- 31 ReLU()
- 32 Linear(in\_features=32, out\_features=num\_classes)
- 33 )

34 Ans: As shown above, the model that is used to recognize handwritten digits is a fully connected model composed of three hidden layers. As with any fully-connected model, each node in a layer is connected to all nodes in the next layer. ReLU is the activation function for the first and the second hidden layers and it serves to introduce nonlinearity into the model.

- 35
- 36 • Loss function: Cross-Entropy Loss

37 Ans: Cross-Entropy is the loss function of choice as it is designed to measure differences between classes. The implementation of Cross-Entropy loss in PyTorch utilizes the sigmoid function which ensures all values fall between 0 and 1 and the negative log loss where loss goes to 0 as input approaches one.

- 38
- 39 • Optimizer: Adam

40 Ans: Adam is an adaptive optimizer that works well in most cases. It utilizes an adaptive learning rate where each weight has its own learning rate and momentum where

learning rates get adjusted based on the gradient direction and time.

```
1 from pathlib import Path
2 from scipy.io import loadmat
3 from sklearn.model_selection import KFold
4 import torch
5 from torch.nn import CrossEntropyLoss
6 from torch.optim import Adam
7 from torch.utils.data import DataLoader
8 from dataset import CustomDataset
9 from model import Model
10 from sklearn.model_selection import train_test_split
11 from utilities import train, test
12
13 ROOT = Path(__file__).parent
14
15
16 def problem_2(
17     , data_path: str = ROOT / "digits.mat", labels_path: str = ROOT / "labels.mat"
18     ,):
19     """Perform K-Fold validation and use those data to train a neural network. The training
20     process utilizes early stopping to prevent overfitting.
21
22     Args:
23         data_path (str, optional): path to the data file. Defaults to ROOT/"digits.mat".
24         labels_path (str, optional): path to the labels file. Defaults to ROOT/"labels.mat".
25
26     """
27     # Hyperparameters
28     n_splits = 15
29     lr = 0.001
30     batch_size = 32
31     num_epochs = 1000
32     patient = 10
33
34     # Load data from .mat files
35     data = loadmat(data_path)["data"].T
36     labels = loadmat(labels_path)["labels"].T.flatten()
37
38     # K-Fold cross validation
39     kf = KFold(n_splits, shuffle=True)
40
41     # Keep track of current fold
42     split = 0
43
44     # Initialize average metrics
45     avg_train_loss = 0
46     avg_train_acc = 0
47     avg_val_loss = 0
48     avg_val_acc = 0
49     avg_test_loss = 0
50     avg_test_acc = 0
51
52     # Perform k-fold cross validation
53     for train_indices, test_indices in kf.split(data):
54         # Split dataset into train, valid, and test set
55         X_train, X_val, y_train, y_val = train_test_split(
56             data[train_indices], labels[train_indices]
```

```
57
58 # Convert numpy array to dataset
59 train_dataset = CustomDataset(X_train, y_train)
60 val_dataset = CustomDataset(X_val, y_val)
61 test_dataset = CustomDataset(X_test, y_test)
62
63 # Load train dataset into dataloader for batch training
64 train_dataloader = DataLoader(train_dataset, batch_size)
65
66 # Define features and classes
67 in_shape = torch.from_numpy(data).size()
68 num_classes = torch.unique(torch.from_numpy(labels)).size()[0]
69
70 # Define model, loss function, and optimizer
71 model = Model(in_shape, num_classes)
72 loss_fn = CrossEntropyLoss()
73 optimizer = Adam(model.parameters(), lr)
74
75 # Initialize per fold metrics
76 train_loss = 0
77 train_acc = 0
78 val_loss = 0
79 val_acc = 0
80 test_loss = 0
81 test_acc = 0
82
83 # Early Stop: store multiple losses
84 es_loss_list = []
85
86 # Train the model
87 for epoch in range(num_epochs):
88
89     # Batch training
90     train_loss, train_acc, model = train(
91         model, loss_fn, optimizer, train_dataloader
92     )
93
94     # Validate the model
95     val_loss, val_acc, model = test(model, loss_fn, val_dataset)
96
97     # Print per epoch metrics
98     print(
99         f"Epoch: {epoch+1}/{num_epochs}",
100         f" train_loss:{train_loss:.5f}",
101         f" train_acc: {train_acc:.3f}",
102         f" val_loss:{val_loss:.5f}",
103         f" val_acc:{val_acc:.3f}",
104         end="\n",
105     )
106
107     # Early stop algorithm
108     # Add latest loss to the end of the list
109     es_loss_list.append(val_loss)
110
111     # Remove first loss if the list is larger than patient
112     if len(es_loss_list) > patient:
113         es_loss_list = es_loss_list[1:]
```

```
114
115     # Check if the list is not decreasing
116     not_decreasing = all(a <= b for a, b in zip(es_loss_list, es_loss_list[1:]))
117
118     # Stop training once val loss list not decreasing
119     if not_decreasing and len(es_loss_list) == patient:
120         break
121
122     # Test the model
123     test_loss, test_acc, model = test(model, loss_fn, test_dataset)
124
125     # Print per fold metrics
126     print(
127         f"Fold: {split+1}/{n_splits},",
128         f" test_loss:{test_loss:.5f},",
129         f" test_acc:{test_acc:.3f},",
130         f" train_loss:{train_loss:.5f},",
131         f" train_acc: {train_acc:.3f},",
132         f" val_loss:{val_loss:.5f},",
133         f" val_acc:{val_acc:.3f}",
134     )
135
136     # Update average metrics
137     avg_val_loss += val_loss / n_splits
138     avg_val_acc += val_acc / n_splits
139     avg_train_loss += train_loss / n_splits
140     avg_train_acc += train_acc / n_splits
141     avg_test_loss += test_loss / n_splits
142     avg_test_acc += test_acc / n_splits
143
144     # Update split
145     split += 1
146
147     # Print average metrics
148     print(
149         "----Average Metrics----\n",
150         f"num_folds: {n_splits}\n",
151         f"avg_test_acc: {avg_test_acc:.3f}\n",
152         f"avg_train_acc: {avg_train_acc:.3f}\n",
153         f"avg_val_acc: {avg_val_acc:.3f}\n",
154         f"avg_test_loss: {avg_test_loss:.5f}\n",
155         f"avg_train_loss: {avg_train_loss:.5f}\n",
156         f"avg_val_loss: {avg_val_loss:.5f}",
157     )
158
159
160 if __name__ == "__main__":
161     problem_2()
```



```
1 import re
2 import torch
3 from torch.nn import Module, CrossEntropyLoss
4 from torch.utils.data import DataLoader, Dataset
5 from torch.optim import Adam
6
7
8 def train(
9     model: Module, loss_fn: CrossEntropyLoss, optimizer: Adam, dataloader: DataLoader
10 ) -> tuple:
11     """Train a given model.
12
13     Args:
14         model (Module): model.
15         loss_fn (CrossEntropyLoss): loss function.
16         optimizer (Adam): optimizer.
17         dataloader (DataLoader): dataloader.
18
19     Returns:
20         tuple: loss, accuracy, model.
21     """
22     # Keep tracking for y and y_pred to calculate final loss and accuracy
23     all_y = None
24     all_y_pred = None
25
26     # Perform batch training
27     for _, (X, y) in enumerate(dataloader):
28
29         # Forward Propagation
30         y_pred = model(X)
31
32         # Calculate loss
33         loss = loss_fn(y_pred, y)
34
35         # Calculate gradients
36         loss.backward()
37
38         # Update weights
39         optimizer.step()
40
41         # Clearn gradients in the optimizer
42         optimizer.zero_grad()
43
44         # Store y and y_pred of this batch
45         with torch.no_grad():
46             all_y = y if all_y == None else torch.cat((all_y, y))
47             all_y_pred = (
48                 y_pred if all_y_pred == None else torch.cat((all_y_pred, y_pred))
49             )
50
51     # Calculate loss and accuracy
52     loss, acc = calculate_loss_accuracy(all_y_pred, all_y, loss_fn)
53
54     return loss, acc, model
55
56
57 def test(model: Module, loss_fn: CrossEntropyLoss, dataset: Dataset) -> tuple:
```

```
58     """Test a given model.
59
60     Args:
61         model (Module): model.
62         loss_fn (CrossEntropyLoss): loss function.
63         dataset (Dataset): dataset.
64
65     Returns:
66         tuple: loss, accuracy, model
67     """
68     # Use no_grad to freeze the model.
69     with torch.no_grad():
70         X = dataset.X
71         y = dataset.y
72
73         # Forward Propagation
74         y_pred = model(X)
75
76         # Calculate loss and accuracy
77         loss, acc = calculate_loss_accuracy(y_pred, y, loss_fn)
78
79     return loss, acc, model
80
81
82 def calculate_loss_accuracy(
83     , y_pred: torch.Tensor, y: torch.tensor, loss_fn: CrossEntropyLoss
84     ,) -> tuple:
85     """Calculate loss and accuracy with given labels and predicted labels.
86
87     Args:
88         y_pred (torch.Tensor): predicted labels.
89         y (torch.tensor): true labels.
90         loss_fn (CrossEntropyLoss): loss function.
91
92     Returns:
93         tuple: loss, accuracy
94     """
95     acc = torch.sum(torch.argmax(y_pred, 1) == y) / y.size()[0]
96     loss = loss_fn(y_pred, y).item()
97     return loss, acc
```

```
1 import torch
2 from torch.nn import Module, Linear
3 from torch.nn.functional import relu
4
5
6 class Model(Module):
7     """An implementation of torch.nn.Module.
8
9     Args:
10         Module (Class): generic pytorch model class.
11     """
12
13     def __init__(self, in_shape: torch.Size, num_classes: int):
14         """Initialize the model
15
16         Args:
17             in_shape (torch.Size): the shape of input.
18             num_classes (int, optional): number of output classes.
19         """
20         super(Model, self).__init__()
21
22         # Parameters
23         self.in_features = torch.prod(torch.tensor(in_shape[1:]))
24         self.num_classes = num_classes
25
26         # Define layers
27         self.fc0 = Linear(self.in_features, 32)
28         self.fc1 = Linear(32, 32)
29         self.fc2 = Linear(32, self.num_classes)
30
31     def forward(self, x: torch.Tensor) -> torch.Tensor:
32         """Feed data through the model.
33
34         Args:
35             x (torch.Tensor): data.
36
37         Returns:
38             torch.Tensor: label.
39         """
40         x = relu(self.fc0(x))
41         x = relu(self.fc1(x))
42         x = self.fc2(x)
43         return x
```

```
1 from torch.utils.data import Dataset
2 import torch
3 import numpy as np
4
5
6 class CustomDataset(Dataset):
7     """An implementation of torch.utils.data.Dataset .
8
9     Args:
10         Dataset (Class): generic pytorch dataset class.
11     """
12
13     def __init__(self, X: np.ndarray, y: np.ndarray):
14         """Initialize the dataset
15
16         Args:
17             X (np.ndarray): data.
18             y (np.ndarray): labels.
19         """
20         super(CustomDataset, self).__init__()
21         self.X = torch.from_numpy(X).float()
22         self.y = torch.from_numpy(y)
23
24     def __getitem__(self, index: int) -> tuple:
25         """Return data and label based on index.
26
27         Args:
28             index (int): index.
29
30         Returns:
31             tuple: data, label
32         """
33         return self.X[index], self.y[index]
34
35     def __len__(self) -> int:
36         """Return dataset length
37
38         Returns:
39             int: length
40         """
41         return self.X.shape[0]
```